


Overestimation bias along training (mean and 10th-90th percentile interval)

black dots: steps where LN and no LN differ (Welch's t-test, $p < 0.05$, uncorrected)

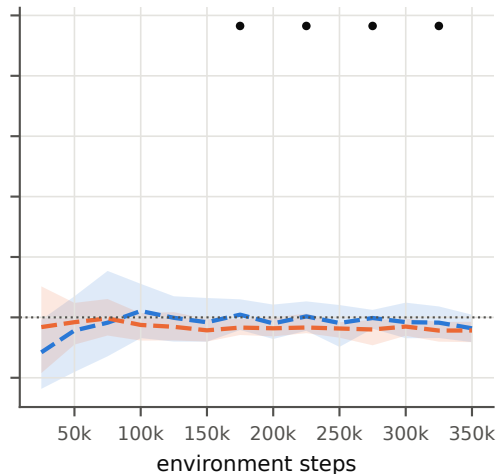
DDPG



LN vs no LN: Welch $p < 0.001$, Bonferroni $p = 0.006$

— no LN (n = 20) — LN (n = 20)

TD3



LN vs no LN: Welch $p < 0.001$, Bonferroni $p = 0.001$

- - no LN (n = 20) - - LN (n = 20)